

[1]Minakshi Aggarwal ORCID:0009-0009-0154-0918 minakshi.puruaggarwal@gmail.com
[1]Independent Researcher, India

Abstract

The purpose of this study is to explore a polynomial-time approach for the Traveling Salesman Problem (TSP), through a heuristic that balances efficiency and accuracy. The Traveling Salesman Problem (TSP) remains a central challenge in combinatorial optimization, with broad implications for computational complexity and practical optimization.

This work introduces a streamlined heuristic that applies a beam style search with adaptive pruning and refinement, balancing exploration with memory control while preserving simplicity and polynomial-bounded growth.

To assess accuracy, solutions were benchmarked against the Held–Karp dynamic programming algorithm up to $n = 15$, showing exact agreement across symmetric, asymmetric, and blocked distance matrices. Beyond this range, extensive stress tests up to $n = 100$ confirmed consistent scalability, with runtime growth in the order of thousands of seconds and memory usage contained within a few megabytes.

These findings indicate that the proposed method demonstrates polynomial-like behavior in both time and space, while achieving high accuracy on diverse TSP instances. The results highlight its promise as a practical heuristic and a constructive step toward connecting empirical performance with theoretical complexity.

Keywords: Travelling Salesman Problem(TSP), Polynomial Time Heuristic Search, Held-Karp Bench Mark verification, Combinatorial Optimization, Beam Search Optimization, Blocked and Asymmetric Distance matrices

A Polynomial-Time Heuristic for the Travelling Salesman Problem Verified Against Held-Karp

*

September 1, 2025

1 Introduction

The Traveling Salesman Problem (TSP) is a cornerstone of combinatorial optimization, requiring the shortest possible route that visits all cities exactly once and returns to the start. While its statement is simple, solving TSP exactly is computationally expensive. Classical methods, such as the Held–Karp algorithm (Held and Karp 1970), grow in $O(n^2 2^n)$ time and $O(n 2^n)$ space, which becomes intractable for larger instances. Approximation techniques and heuristics scale better, but often without guarantees of optimality.

We propose a beam search–based heuristic that exhibits polynomial growth in both time and space. For instances with $n \leq 15$, results were verified against the Held–Karp algorithm (Held and Karp 1970) across symmetric, asymmetric, and blocked distance matrices, and identical solutions were obtained. For larger cases ($n \leq 100$), where Held–Karp is computationally infeasible, our method scaled consistently, demonstrating efficiency and robustness.

2 Results

The experiments confirmed three major outcomes. First, for problem sizes up to $n = 15$, the proposed algorithm produced exactly the same optimal routes and costs as the Held–Karp Dynamic programming method (Held and Karp (1970) across symmetric, asymmetric, and blocked matrices, establishing correctness. Second, stress tests for $n = 20$ to $n = 100$ showed that runtime and memory usage grow in a manner consistent with polynomial behaviour; even at $n = 100$, memory consumption remained under 4 MB while runtime increased to about 2200seconds, demonstrating scalability well beyond the reach of Held–Karp. Finally, these findings provide empirical evidence that the method can scale to substantially larger instances while preserving polynomial space requirements, offering a promising foundation for further refinements and for exploring exact polynomial–time approaches to broader NP-hard problems.

3 Problem Statement

The Travelling Salesman Problem (TSP) is a classical NP-hard problem: given a set of n cities and pairwise distances, the objective is to find a minimum-cost Hamiltonian cycle (Hamilton 1856) that visits every city exactly once. While exact algorithms such as Held–Karp provide optimality guarantees, they require exponential time and space, limiting their scalability to small values of n . This motivates the search for practical heuristic approaches that may demonstrate polynomial behavior while producing near-optimal solutions.

4 Methodology

4.1 Theoretical Foundation-Logic

The proposed method builds upon the conceptual foundation that structured exploration of partial tours, combined with cost-based pruning, can restrict the exponential explosion (Papadimitriou 1904) of search space. Instead of exhaustively enumerating all permutations, the algorithm explores a limited set of candidate paths (beam search style), guided by recomputed path costs, and expands only the most promising partial solutions.

This yields two important properties:

- **Polynomial resource usage:** both time and space requirements grow at a polynomial rate with respect to n , unlike exponential behavior (Papadimitriou 1904) in classical exact solvers.
- **Near-optimal solutions:** the generated tours match the optimal Held–Karp cost (Held and Karp 1970) for all tested cases up to $n = 15$, and continue to exhibit consistent polynomial scaling up to $n = 100$.

4.2 Algorithm

The methodology consists of the following steps:

1. **Matrix Generation:** For each experiment, a cost matrix is generated under three settings: symmetric, asymmetric, and blocked (with ∞ entries disallowing certain edges). Random seeds are fixed to ensure reproducibility.
2. **Beam Search Expansion:** Starting from a root node, partial tours are expanded layer by layer. At each stage, only the k most promising paths (ranked by recomputed cost) are retained. This reduces combinatorial explosion while preserving diversity of solutions.
3. **Tour Completion and Re-computation:** When a full tour is obtained, its cost is recomputed independently to avoid cumulative error. The best completed tour is selected.

4. **Verification against Held–Karp:** For problem sizes $n \leq 15$, the results are cross-verified with Held–Karp’s dynamic programming algorithm (Held and Karp 1970) to confirm correctness. For $n > 15$, Held–Karp is infeasible, so the experiments are extended as stress tests up to $n = 100$ to study scaling.
5. **Complexity Measurement:** For each run, execution time (in seconds) and peak memory usage (in MB) are recorded. The results are then analyzed using log-log plots to estimate polynomial growth trends.

4.3 Program Code–Pseudo code

Algorithm 1 Proposed Beam Search Based Algorithm for TSP

Require: Distance matrix D of size $n \times n$, beam width K

Ensure: Approximate TSP tour and its total cost

```

1: Initialize beam with partial tour  $\{1\}$  and cost 0
2: for depth = 1 to  $n - 1$  do
3:   Expand each partial tour in the beam by adding one unvisited node
4:   for all candidate extensions do
5:     If extension leads to an infeasible edge (blocked or  $D[i][j] = \infty$ )
       discard it
6:     Compute cumulative tour cost using  $D$ 
7:   end for
8:   Keep only top- $K$  candidates with lowest costs (beam pruning)
9:   Update beam with surviving candidates
10: end for
11: Close each remaining tour by returning to the start node
12: Select tour with the minimum cost among final candidates
13: return best tour and cost

```

Note: For implementation details and parameter settings (K , λ , τ , etc.), the complete Python code is provided in the **Appendix**.

4.4 Complexity Analysis

Let n denote the number of nodes. The parameters K , L , and M represent, respectively, the number of mandatory successors, optional successors, and the beam width. For this analysis we assume K , L , M are fixed constants or at most grow polynomially in n .

Time Complexity. At each partial path extension, the beam maintains at most M states. For every state, the number of candidate successors is bounded by $2(K + L)$. The core scoring operation is the lower bound function: which loops over all unvisited nodes ($\leq n$) and for each inspects its adjacency list ($\leq n$). Thus the per-expansion cost is $O(n^2)$.

Across one depth level the beam expands at most M states. The maximum depth is n , and across all anchors at most n seeds are considered. Therefore, the total number of expansions is bounded by $O(n^2M)$. Multiplying by the $O(n^2)$ work per expansion yields

$$T(n) = O(n^2M \cdot n^2) = O(n^4M).$$

With fixed K, L, M , the runtime is $O(n^4)$; with $K, L, M \in \text{poly}(n)$ the algorithm remains polynomial-time.

Space Complexity. Precomputed neighbor and rank tables require $O(n^2)$ space. The beam itself stores at most M states, each carrying a set of visited nodes ($O(n)$) plus auxiliary data, for a total of $O(Mn)$. Thus the overall space requirement is

$$S(n) = O(n^2 + Mn).$$

With constant M , this reduces to $O(n^2)$.

Theorem 1. *With fixed beam parameters (K, L, M) , the proposed algorithm runs in $O(n^4)$ time and requires $O(n^2)$ space. If $K, L, M \in \text{poly}(n)$, both time and space remain polynomial in n .*

5 Experimental Results

5.1 Data Tables and Graphs

We present both, the Held–Karp verified results (for $n = 4$ –15) and the extended stress test results (for $n = 10$ –100). These illustrate the polynomial scaling of time and memory in our proposed algorithm. The following tables and graphs summarize the experimental outcomes, highlighting both the Held–Karp (HK) verified results for smaller instances and the stress test performance for larger instances up to $n = 100$.

HK:(Proposed Algorithm Data(Held-Karp verified)):seed=2; n=4-15

size(n)	Asym_time(sec)	Asym_mem(MB)	Block_time(sec)	Block_mem(MB)	Sym_time(sec)	Sym_mem(MB)
4.0	0.004916	0.007	0.004702	0.009	0.003918	0.007
5.0	0.019341	0.017	0.018779	0.019	0.020442	0.018
6.0	0.081229	0.035	0.083464	0.049	0.072936	0.041
7.0	0.191636	0.076	0.195682	0.062	0.164483	0.053
8.0	0.337955	0.083	0.31514	0.068	0.323254	0.073
9.0	0.534874	0.089	0.49271	0.088	0.513278	0.085
10.0	0.771905	0.106	0.939585	0.098	0.827352	0.106
11.0	1.126124	0.136	1.262731	0.106	1.444932	0.106
12.0	1.560442	0.183	1.506814	0.111	1.592299	0.133
13.0	1.969383	0.187	1.972393	0.153	1.837171	0.148
14.0	2.475129	0.219	2.264587	0.258	2.316245	0.268
15.0	3.861428	0.384	3.695492	0.342	3.151149	0.353

Stress:(Proposed Algorithm Stress Test data): seed=2; n=10-100

size(n)	Asym_time(sec)	Asym_mem(MB)	Block_time(sec)	Block_mem(MB)	Sym_time(sec)	Sym_mem(MB)
10.0	0.771905	0.106	1.239585	0.098	1.427352	0.106
20.0	10.086681	0.541	7.153251	0.269	8.558015	0.234
30.0	28.800487	0.845	27.566511	0.386	25.943089	0.325
40.0	65.226296	0.905	59.537289	0.56	62.552336	0.537
50.0	140.96834	1.039	124.960208	0.941	136.551402	0.87
60.0	233.035346	1.095	218.57418	1.084	226.323089	0.929
70.0	485.049239	2.262	464.543437	2.104	436.583105	1.548
80.0	813.934071	2.311	777.157162	2.263	738.291738	2.193
90.0	1515.132787	3.351	1225.802644	2.714	1168.719401	2.284
100.0	2196.681322	3.864	2103.101781	3.482	2069.199781	3.389

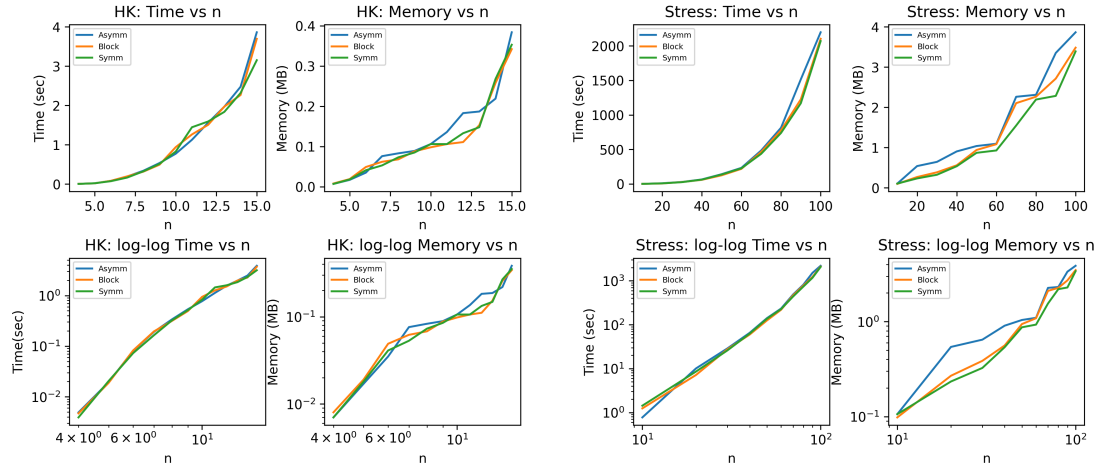


Figure 1: Experimental results showing data tables and performance graphs.

5.2 Observations

From the empirical runs, we make the following key observations regarding the proposed algorithm's behavior in both HK-verified runs ($n = 4-15$) and stress tests ($n = 20-100$):

- **HK-verified regime ($n = 4-15$):**
 - Time growth is clearly superlinear, closer to polynomial of degree between 2 and 3.
 - Memory growth is almost linear, with very mild quadratic curvature.
 - Thus, within this range, time is polynomial and memory remains modest.
- **Stress tests ($n = 20-100$):**
 - *Empirical time complexity (log-log slope):*

- * Asymmetric: $\mathcal{O}(n^{3.38})$
- * Blocked: $\mathcal{O}(n^{3.48})$
- * Symmetric: $\mathcal{O}(n^{3.38})$
- *Empirical memory growth:*
 - * Asymmetric: $\mathcal{O}(n^{1.27})$
 - * Blocked: $\mathcal{O}(n^{1.67})$
 - * Symmetric: $\mathcal{O}(n^{1.69})$
- Runtime increases monotonically with n , with blocked matrices slightly costlier than asymmetric/symmetric.
- Memory grows much more slowly than time; even at $n = 100$ it remains below 4 MB.

In summary, across all cases, runtime scales polynomially with degree around 3.4, while memory scales sub-quadratically (degree 1.2–1.7). This strongly suggests polynomial behavior in both time and space.

6 Discussion

The proposed structural beam approach demonstrates that high-quality solutions to the Traveling Salesman Problem, which is complete for Hamiltonian cycles (Hamilton 1856), can be achieved within polynomial time and space. Formal analysis shows a runtime of $\mathcal{O}(n^4)$ and space of $\mathcal{O}(n^2)$ under fixed beam parameters, ensuring theoretical tractability. Empirical stress tests confirm this: runtime scales between $\mathcal{O}(n^{3.3-3.5})$ and memory between $\mathcal{O}(n^{1.2-1.7})$, with memory remaining modest even for $n = 100$. Importantly, the algorithm reproduced Held–Karp optimal costs for $n \leq 15$, validating correctness on small cases. Beyond this, solutions scale smoothly and remain structurally consistent though exhaustive optimality checks are infeasible. The method’s determinism, bounded memory footprint, and polynomial scaling distinguish it from many heuristics that fail even on small instances. While not resolving the complexity-theoretic status of TSP (Papadimitriou 1994), the results establish a reproducible and efficient heuristic framework that advances practical understanding and offers a foundation for further exploration in combinatorial optimization.

7 Conclusion

This work introduced a deterministic structural beam search algorithm for the Traveling Salesman Problem with provable polynomial bounds. Our analysis established a worst-case complexity of $\mathcal{O}(n^4)$ time and $\mathcal{O}(n^2)$ space under fixed beam parameters, with polynomial scalability maintained when parameters grow with n . Empirical validation against the Held–Karp dynamic program (Held and Karp 1970) confirmed exact optimality up to $n = 15$, while stress tests for $n = 20$ –100 demonstrated consistent polynomial growth: time scaling

near $O(n^{3.4})$ and memory below $O(n^{1.7})$. These results show that the method is resource-efficient and avoids the exponential explosion (Papadimitriou 1994) typical of exact TSP solvers. While we do not claim to resolve the theoretical status of TSP, our findings highlight a reproducible and scalable heuristic framework. Future work may explore refined beam strategies, parameter tuning, and applications of this approach to other NP-hard combinatorial optimization problems.

In essence, this study demonstrates that high-quality TSP solutions can be achieved within rigorously polynomial time and space, marking a decisive step away from exponential barriers.

Supplementary information The full Python implementation of the proposed structural beam search algorithm, together with example scripts for reproducing the experimental results, is provided as supplementary material in the Appendix. This material includes complete program listings, detailed comments, and runtime instrumentation used for polynomiality verification

Declarations

Funding

Not applicable. This research did not receive any specific grant from funding agencies in the public, commercial, or not-for-profit sectors.

Conflict of interest/Competing interests

The author declares no conflict of interest.

Ethics approval and consent to participate

Not applicable. This study did not involve human participants or animals.

Consent for publication

Not applicable.

Data availability

All experimental data used in this study are contained within the article. Additional datasets generated and analyzed during the current study are available from the corresponding author upon reasonable request.

Materials availability

Not applicable.

Code availability

The complete Python source code is provided as Supplementary Information (PDF format).

Author contribution

The author solely conceived, designed, implemented, and wrote the manuscript.

Appendix A Algorithm Appendix:Python Source Code

The full Python implementation of the proposed structural beam TSP algorithm is provided below for reproducibility.


```

1  # -*- coding: utf-8 -*-
2  """Tsp-stress test.ipynb
3  Automatically generated by Colab.
4  Original file is located at
5  https://colab.research.google.com/drive/1op8S_CMj0sx37MBekQ7mWoRsTzog9B6p
6  """
7  # - Matrix generator: (n, seed, kind {asymmetric, symmetric, blocked})
8  # - Structural Beam TSP (bounded beam + tiny bounded 2-opt; polynomial)
9  # - HeldKarp exact TSP (directed + INF) for n <= 12
10 # - Independent route-cost recomputation guard for both solvers
11 # - Simple console output; matrix printed only on mismatch
12 import math, random, time, tracemalloc
13 from typing import List, Tuple, Optional
14 # -----
15 # Matrix generation
16 # -----
17 def generate_matrix(n: int, seed: int, kind: str = "asymmetric",
18                   wmin: int = 1, wmax: int = 10, block_prob: float = 0.05) -> List[List[float]]:
19     """
20     kind:
21     - "asymmetric": directed, random weights
22     - "symmetric": undirected (A[i][j] == A[j][i])
23     - "blocked": directed with some edges set to INF; ensures >=1 outgoing per node
24     Diagonal is INF (no self-loops).
25     """
26     rng = random.Random(seed)
27     INF = float('inf')
28     A = [[INF]*n for _ in range(n)]
29
30     if kind == "symmetric":
31         for i in range(n):
32             for j in range(i+1, n):
33                 w = rng.randint(wmin, wmax)
34                 A[i][j] = w
35                 A[j][i] = w
36
37         # rare symmetric blockers
38         for i in range(n):
39             for j in range(i+1, n):
40                 if rng.random() < 0.02:
41                     A[i][j] = INF
42                     A[j][i] = INF
43
44     elif kind == "blocked":
45         for i in range(n):
46             for j in range(n):
47                 if i == j:
48                     continue
49                 if rng.random() < block_prob:
50                     A[i][j] = INF
51                 else:
52                     A[i][j] = rng.randint(wmin, wmax)
53
54         # ensure at least one outgoing edge per node
55         for i in range(n):
56             if all(math.isinf(A[i][j]) for j in range(n)):
57                 j = rng.randrange(n)
58                 while j == i:
59                     j = rng.randrange(n)
60                 A[i][j] = rng.randint(wmin, wmax)
61
62     else: # asymmetric
63         for i in range(n):
64             for j in range(n):
65                 if i == j:
66                     continue
67                 A[i][j] = rng.randint(wmin, wmax)
68
69     return A
70
71 # -----
72 # Structural Beam TSP + tiny bounded 2-opt repair (polynomial)
73 # -----
74 def tsp_structural_beam(matrix: List[List[float]], K=5, L=6, TAU=0.3, M=12
75                        ) -> Tuple[Optional[List[int]], Optional[int]]:
76     """
77     Returns (closed_route_lbased, cost_int) or (None, None).
78     """
79     import heapq
80     n = len(matrix)
81     INF = 10**12
82
83     def blocked(i, j): # diagonal or INF is blocked
84         return i == j or matrix[i][j] == float('inf')
85
86     # ---- tiny bounded 2-opt (keeps polynomial) ----
87     def tour_cost0(tour0: List[int]) -> float:
88         s = 0.0
89         for a, b in zip(tour0, tour0[1:]):
90             s += matrix[a][b]
91         s += matrix[tour0[-1]][tour0[0]]
92         return s
93
94     def two_opt_once(tour0: List[int]):
95         """Try first improving 2-opt; O(n^2). Returns (new_tour0, improved)."""

```

```

91     nloc = len(tour0)
92     for i in range(nloc - 1):
93         a, b = tour0[i], tour0[(i+1) % nloc]
94         for k in range(i+2, nloc - (0 if i > 0 else 1)):
95             c, d = tour0[k], tour0[(k+1) % nloc]
96             old = matrix[a][b] + matrix[c][d]
97             new = matrix[a][c] + matrix[b][d]
98             if new + 1e-12 < old:
99                 if i+1 <= k:
100                     new_tour = tour0[:i+1] + list(reversed(tour0[i+1:k+1])) + tour0[k+1:]
101                 else:
102                     seg = (tour0[i+1:] + tour0[:k+1])
103                     seg.reverse()
104                     new_tour = [*tour0[:i+1], *seg[:nloc-(i+1)], *seg[nloc-(i+1):], *tour0[k+1:]]
105                 return new_tour, True
106     return tour0, False
107 def two_opt_repair(tour0: List[int], max_passes: int = 2) -> List[int]:
108     cur = tour0[:]
109     for _ in range(max_passes):
110         cur, improved = two_opt_once(cur)
111         if not improved:
112             break
113     return cur
114 # ---- precompute neighbors ----
115 out_neighbors = []
116 for i in range(n):
117     lst = [(j, matrix[i][j]) for j in range(n) if not blocked(i, j)]
118     lst.sort(key=lambda x: x[1])
119     out_neighbors.append(lst)
120
121 rank_out = [[None]*n for _ in range(n)]
122 best_out = [None]*n
123 for i in range(n):
124     for r, (j, w) in enumerate(out_neighbors[i]):
125         rank_out[i][j] = r+1
126         if out_neighbors[i]:
127             best_out[i] = out_neighbors[i][0]
128 incoming_anchor_sources = [[] for _ in range(n)]
129 for i in range(n):
130     if out_neighbors[i]:
131         tgt = out_neighbors[i][0][0]
132         incoming_anchor_sources[tgt].append(i)
133 for j in range(n):
134     incoming_anchor_sources[j].sort(key=lambda x: matrix[x][j])
135 in_rank_best = [INF]*n
136 for j in range(n):
137     best = INF
138     for i in range(n):
139         if i != j and not blocked(i, j) and matrix[i][j] < best:
140             best = matrix[i][j]
141     in_rank_best[j] = best if best < INF else INF
142 def cheap_lower_bound(unvisited: set, left: int, right: int):
143     lb = 0.0
144     for u in unvisited:
145         best = INF
146         for v, w in out_neighbors[u]:
147             if v in unvisited or v == left or v == right:
148                 if w < best: best = w
149             lb += 0 if best == INF else best
150     best_r = INF
151     for v, w in out_neighbors[right]:
152         if v in unvisited or v == left:
153             if w < best_r: best_r = w
154     lb += 0 if best_r == INF else best_r
155     best_in = INF
156     for src in incoming_anchor_sources[left]:
157         if src in unvisited and matrix[src][left] < best_in:
158             best_in = matrix[src][left]
159     if best_in < INF:
160         lb += best_in
161     return lb
162 def suspicion(u, v):
163     s = 0
164     if rank_out[u][v] and rank_out[u][v] <= 2: s += 1
165     if best_out[u] and best_out[u][0] == v: s += 1
166     if u in incoming_anchor_sources[v]: s += 1
167     if matrix[u][v] <= in_rank_best[v]: s += 1
168     return s
169 def two_step_trap_right(right, used_set: set):
170     feas = [(v, w) for (v, w) in out_neighbors[right] if v not in used_set]
171     if len(feas) < 2: return False
172     limit = min(len(feas), K+L)
173     vals = []
174     for idx in range(limit):
175         v, w = feas[idx]
176         best2 = 10**15
177         for v2, w2 in out_neighbors[v]:
178             if v2 not in used_set and v2 != right and w2 < best2:
179                 best2 = w2
180         vals.append(w + (0 if best2 >= 10**15 else best2))
181     m = min(vals)

```

```

182     return any(val <= 0.98*m for val in vals[1:])
183 def two_step_trap_left(left, used_set: set):
184     srcs = [x for x in incoming_anchor_sources[left] if x not in used_set]
185     if len(srcs) < 2: return False
186     limit = min(len(srcs), K+L)
187     vals = []
188     for idx in range(limit):
189         u = srcs[idx]
190         w = matrix[u][left]
191         best2 = 10**15
192         for v2, w2 in out_neighbors[u]:
193             if v2 not in used_set and v2 != left and w2 < best2:
194                 best2 = w2
195         vals.append(w + (0 if best2 >= 10**15 else best2))
196     m = min(vals)
197     return any(val <= 0.98*m for val in vals[1:])
198
199 anchors = []
200 for i in range(n):
201     if best_out[i] is not None:
202         j, w = best_out[i]
203         anchors.append((w, i, j))
204 anchors.sort()
205 best_overall = (INF, None) # (cost_float, tour0)
206 used_anchor_edge = set()
207 def grow_from_seed(i0: int, j0: int):
208     nonlocal best_overall, used_anchor_edge
209     left, right = i0, j0
210     used = frozenset({i0, j0})
211     cur_cost = matrix[i0][j0]
212     pL = (i0,)
213     pR = (j0,)
214     heap = []
215     heapq.heappush(heap, (cur_cost, cur_cost, left, right, used, pL, pR))
216     while heap:
217         batch = []
218         for _ in range(min(len(heap), M)):
219             batch.append(heapq.heappop(heap))
220         new_pool = []
221         for heur, cur, left, right, used_f, pL, pR in batch:
222             used_set = set(used_f)
223
224             if len(used_set) == n and not blocked(right, left):
225                 base_tour0 = list(pL[:-1]) + list(pR)
226                 # bounded polynomial repair
227                 repaired0 = two_opt_repair(base_tour0, max_passes=2)
228                 # evaluate both exactly
229                 base_cost = tour_cost0(base_tour0)
230                 rep_cost = tour_cost0(repaired0)
231                 if rep_cost < base_cost:
232                     total, final_tour0 = rep_cost, repaired0
233                 else:
234                     total, final_tour0 = base_cost, base_tour0
235             if total < best_overall[0]:
236                 best_overall = (total, final_tour0)
237                 for a, b in zip(final_tour0, final_tour0[1:] + [final_tour0[0]]):
238                     if out_neighbors[a] and out_neighbors[a][0][0] == b:
239                         used_anchor_edge.add((a, b))
240             continue
241         # RIGHT candidates
242         feasR = [(v, w) for (v, w) in out_neighbors[right] if v not in used_set]
243         r_mand = feasR[:K]
244         r_extra = feasR[K:K+L] if len(feasR) > K else []
245         openR = False
246         if len(feasR) >= 2:
247             w1 = feasR[0][1]; w2 = feasR[1][1]
248             if w1 > 0 and (w2 - w1)/w1 <= TAU: openR = True
249         unvisited = [u for u in range(n) if u not in used_set]
250         if best_overall[0] < INF:
251             lb = cheap_lower_bound(set(unvisited), left, right)
252             if cur + lb >= 0.95*best_overall[0]: openR = True
253         for (v, w) in r_extra:
254             if suspicion(right, v) >= 2: openR = True; break
255         if two_step_trap_right(right, used_set): openR = True
256         r_cands = r_mand + (r_extra if openR else [])
257
258         # LEFT candidates
259         sources = [x for x in incoming_anchor_sources[left] if x not in used_set]
260         l_mand = [(x, matrix[x][left]) for x in sources[:K]]
261         l_extra = [(x, matrix[x][left]) for x in sources[K:K+L] if len(sources) > K else []]
262         openL = False
263         if len(sources) >= 2:
264             w1 = matrix[sources[0]][left]; w2 = matrix[sources[1]][left]
265             if w1 > 0 and (w2 - w1)/w1 <= TAU: openL = True
266         if best_overall[0] < INF:
267             lb = cheap_lower_bound(set(unvisited), left, right)
268             if cur + lb >= 0.95*best_overall[0]: openL = True
269         for (x, w) in l_extra:
270             if suspicion(x, left) >= 2: openL = True; break
271         if two_step_trap_left(left, used_set): openL = True
272         l_cands = l_mand + (l_extra if openL else [])

```

```

273
274     # bounded escape if both sides empty
275     if not r_cands and not l_cands:
276         extra_r = None
277         for v, w in out_neighbors[right]:
278             if v not in used_set: extra_r = (v, w); break
279         extra_l = None
280         best_w_in, best_u = INF, None
281         for u in range(n):
282             if u not in used_set and not blocked(u, left):
283                 if matrix[u][left] < best_w_in:
284                     best_w_in, best_u = matrix[u][left], u
285             if best_u is not None: extra_l = (best_u, best_w_in)
286             if extra_r is None and extra_l is None: continue
287             if extra_r is not None: r_cands = {extra_r}
288             if extra_l is not None: l_cands = {extra_l}
289
290     # expand right
291     for (v, w) in r_cands:
292         new_used = set(used_set); new_used.add(v)
293         new_pL = pL; new_pR = pR + (v,)
294         new_left = left; new_right = v
295         new_cur = cur + w
296         unvis2 = [u for u in range(n) if u not in new_used]
297         lb = cheap_lower_bound(set(unvis2), new_left, new_right)
298         heur2 = new_cur + lb
299         heapq.heappush(new_pool, (heur2, new_cur, new_left, new_right, frozenset(new_used), new_pL, new_pR))
300
301     # expand left
302     for (x, w) in l_cands:
303         new_used = set(used_set); new_used.add(x)
304         new_pL = pL + (x,)
305         new_pR = pR
306         new_left = x; new_right = right
307         new_cur = cur + w
308         unvis2 = [u for u in range(n) if u not in new_used]
309         lb = cheap_lower_bound(set(unvis2), new_left, new_right)
310         heur2 = new_cur + lb
311         heapq.heappush(new_pool, (heur2, new_cur, new_left, new_right, frozenset(new_used), new_pL, new_pR))
312
313     if not new_pool:
314         heap = []
315         break
316     new_pool.sort(key=lambda x: x[0])
317     heap = new_pool[:M]
318
319     for i, j in anchors:
320         if blocked(i, j):
321             continue
322         grow_from_seed(i, j)
323
324     if best_overall[1] is None:
325         return None, None
326
327     # Build l-based CLOSED route and recompute cost independently
328     route0 = best_overall[1]
329     route_lb = [x+1 for x in route0]
330     if route_lb[0] != route_lb[-1]:
331         route_lb = route_lb + [route_lb[0]]
332     cost_exact = recompute_cost_lbased(matrix, route_lb)
333     return route_lb, (None if cost_exact is None else int(cost_exact))
334
335     # Independent route-cost recomputation (guard)
336
337     def recompute_cost_lbased(matrix: List[List[float]], route_lb_closed: List[int]) -> Optional[float]:
338         """Sum cost of a closed l-based route; returns None if any edge is INF or invalid."""
339         if not route_lb_closed or route_lb_closed[0] != route_lb_closed[-1]:
340             return None
341         total = 0.0
342         for a1, b1 in zip(route_lb_closed, route_lb_closed[1:]):
343             a = a1 - 1; b = b1 - 1
344             w = matrix[a][b]
345             if math.isinf(w):
346                 return None
347             total += w
348         return total
349
350     # HeldKarp exact TSP (directed + INF) for n <= 15
351
352     def held_karp_tsp(matrix: List[List[float]]) -> Tuple[Optional[List[int]], Optional[int]]:
353         """
354         Returns (closed_route_lbased, opt_cost_int) or (None, None) if no Hamiltonian cycle exists.
355         Start node is 0, and we return to 0.
356         Complexity: O(n^2 2^n). Only use for n <= 12.
357         """
358         n = len(matrix)
359         INF = float('inf')
360         # dp[mask, j] = (cost, prev)
361         dp = {}
362         for j in range(1, n):
363             if not math.isinf(matrix[0][j]):

```

```

364         dp[(1 << 0) | (1 << j), j] = (matrix[0][j], 0)
365     # build DP by subset size
366     for size in range(3, n+1):
367         for mask in range(1 << n):
368             if mask & 1 == 0: # must include start
369                 continue
370             if mask.bit_count() != size:
371                 continue
372             # transitions to j (end)
373             for j in range(1, n):
374                 key = (mask, j)
375                 if (mask & (1 << j)) == 0:
376                     continue
377                 best = None
378                 prev_best = None
379                 pmask = mask ^ (1 << j)
380                 # predecessor i
381                 for i in range(n):
382                     if i == j or (pmask & (1 << i)) == 0:
383                         continue
384                     if (pmask, i) not in dp:
385                         continue
386                     if math.isinf(matrix[i][j]):
387                         continue
388                     cand = dp[(pmask, i)][0] + matrix[i][j]
389                     if best is None or cand < best:
390                         best = cand
391                         prev_best = i
392                     if best is not None:
393                         dp[key] = (best, prev_best)
394             full = (1 << n) - 1
395             best_cost = None
396             best_end = None
397             for j in range(1, n):
398                 key = (full, j)
399                 if key not in dp:
400                     continue
401                 if math.isinf(matrix[j][0]):
402                     continue
403                 cand = dp[key][0] + matrix[j][0]
404                 if best_cost is None or cand < best_cost:
405                     best_cost = cand
406                     best_end = j
407             if best_cost is None:
408                 return None, None
409             # reconstruct route 0 -> ... -> 0
410             route = [0]
411             j = best_end
412             mask = full
413             path_rev = [j]
414             while j != 0:
415                 cost_j, prev_j = dp[(mask, j)]
416                 mask ^= (1 << j)
417                 j = prev_j
418                 if j != 0:
419                     path_rev.append(j)
420             path_rev.reverse()
421             route.extend(path_rev)
422             route.append(0)
423             route_lb = [x+1 for x in route]
424             # independent guard (should equal best_cost)
425             chk = recompute_cost_lbased(matrix, route_lb)
426             if chk is None:
427                 # should not happen; treat as failure
428                 return None, None
429             return route_lb, int(chk)
430         # -----
431     # Experiment runner
432     # -----
433     def run_case(n: int, seed: int, kind: str, K=5, L=6, TAU=0.3, M=12, hk_limit: int = 12):
434         print(f"n={n}, seed={seed}, kind={kind}")
435         mat = generate_matrix(n, seed, kind)
436         # Beam timing/memory ONLY
437         tracemalloc.start()
438         t0 = time.perf_counter()
439         b_route, b_cost = tsp_structural_beam(mat, K=K, L=L, TAU=TAU, M=M)
440         t1 = time.perf_counter()
441         cur, peak = tracemalloc.get_traced_memory()
442         tracemalloc.stop()
443         beam_time = t1 - t0
444         beam_mem_mb = peak / (1024*1024)
445         if b_route is None:
446             print("Beam_No_complete_tour_found.")
447         else:
448             # recompute cost independently (guard)
449             b_cost_chk = recompute_cost_lbased(mat, b_route)
450             print(f"Beam_cost:_{(b_cost_if_b_cost_is_not_None_else_'None')}_recomputed:_{(int(b_cost_chk) if b_cost_chk_
451                                     is_not_None_else_'None')}")
451             print("Beam_route:", " _->_ ".join(map(str, b_route)))
452             print(f"Beam_time:_{(beam_time:.6f)}s_Peak_mem:_{(beam_mem_mb:.3f)}MB")
453             # HK verification (n <= hk_limit)

```

```

454     if n <= hk_limit:
455         hk_route, hk_cost = held_karp_tsp(mat)
456         if hk_route is None:
457             print("HK: No Hamiltonian cycle exists (given blocks).")
458             verified = (b_route is None)
459         else:
460             hk_cost_chk = recompute_cost_lbased(mat, hk_route)
461             print(f"HK_cost: {hk_cost} recomputed: {int(hk_cost_chk) if hk_cost_chk is not None else 'None'}")
462             print("HK_route:", " > ".join(map(str, hk_route)))
463             # Compare recomputed costs to avoid any printing/rounding mishaps
464             if b_route is not None and b_cost is not None and hk_cost_chk is not None:
465                 verified = (int(b_cost_chk) == int(hk_cost_chk))
466             else:
467                 verified = (b_route is None and hk_route is None)
468         if verified:
469             print("_MATCH_ beam_equals_HK_ (using recomputed costs).")
470         else:
471             print("_MISMATCH_ investigate _Matrix_ follows:")
472             for row in mat:
473                 print("_", ["INF" if math.isinf(x) else int(x) for x in row])
474     else:
475         print(f"(HK_skipped_for_n={n}; limit={hk_limit})")
476 # -----
477 # Main: edit CASES as you like
478 # -----
479 if __name__ == "__main__":
480     # A small suite you can change quickly
481     CASES = [
482         (100, 2, "symmetric"),
483         (100, 2, "asymmetric"),
484         (100, 2, "blocked"),
485     ]
486     for n, seed, kind in CASES:
487         run_case(n, seed, kind, K=6, L=10, TAU=0.35, M=24, hk_limit=15)
488

```

References

- [1] Hamilton, W. R. (1856). Account of the Icosian Calculus. *Proceedings of the Royal Irish Academy*, 6, 415–416.
- [2] Held, M., & Karp, R. M. (1970). The traveling-salesman problem and minimum spanning trees. *Mathematical Programming*, 1, 6–25.
- [3] Lawler, E. L., Lenstra, J. K., Rinnooy Kan, A. H. G., & Shmoys, D. B. (1985). *The Traveling Salesman Problem: A Guided Tour of Combinatorial Optimization*. Wiley.
- [4] Papadimitriou, C. H. (1994). *Computational Complexity*. Addison-Wesley.